

Advanced Database Labs

ADB Labs

Amira Nasreldeen, Aseel, Salma

© 2025 Amira – Aseel – Salma

Table of contents

1. Advanced Database Labs	3
1.1 What You Will Learn	3
1.2 Tutorial Support video	3
1.3 Labs Overview	3
1.4 What You Should Do	3
2. General Instructions	4
2.1 Please follow these instructions carefully:	4
3. Setup environment	5
3.1 Prerequisites: Install Microsoft Visual C++ 2019 Redistributable	6
3.2 MySQL	6
3.3 Python programming language	12
3.4 Alternative way to practice:	14
3.5 Assignment: Install MySQL, Create a Database, and Query It	14
4. Transaction Management and SQL Implementation	15
4.1 Understanding SQL Transactions	15
4.2 Assignment: Simulate The Unrepeatable Read Problem	25
5. Concurrency Control Techniques	27
5.1 InnoDB Storage Engine	27
5.2 Isolation Level on MySQL	27
5.3 InnoDB-Supported Concurrency Control Algorithms	28
5.4 Advanced work	33
5.5 Assignment (Simulate Deadlock Scenario)	33

1. Advanced Database Labs

Welcome to the Advanced Database Lab course. This site provides hands-on labs to help you learn and simulate key concepts in **concurrency control** and **database recovery techniques**.

1.1 What You Will Learn

- How to install and use **MySQL**, **Python**, and **Visual Studio Code**
- An understanding of **transaction management** and **SQL implementation**
- Understanding concurrency control mechanisms (**2PL**, **Timestamp ordering**)
- Understanding Database recovery techniques (**ARIES**)
- An introduction to **database security** practices:
 - **Role-Based Access Control (RBAC)**
 - **Data encryption**

1.2 Tutorial Support video

Many labs include **video walkthroughs** that demonstrate how to solve key theoretical examples step by step.

- Transaction schedules and conflicts
- Concurrency Control mechanisms
- Database recovery Techniques

Be sure to watch them alongside the written material for deeper understanding.

1.3 Labs Overview

Each lab builds on the previous one. Please complete them in sequence:

1. **Lab 0:** Environment Setup
2. **Lab 1:** Basics of Transaction Management and Its Properties
3. **Lab 2:** Concurrency Control Mechanisms
4. **Lab 3:** Database Recovery Techniques (ARIES)
5. **Lab 4:** Database Security Measures



Note

You can download a single PDF that contains all labs. It will be updated each time a new lab is uploaded.

[Download"](#)

1.4 What You Should Do



Each lab includes an assignment that you are expected to complete as part of the learning process. **Please submit the assignments on time.**

2. General Instructions

2.1 Please follow these instructions carefully:

2.1.1 File Submission

- At the top of your document, please include the following information:
- Full Name
- Department
- Index Number
- Submit your assignment as a single PDF file.
- The PDF must include all required screenshots (code and output).

2.1.2 Lab Title

Write the lab title as "Lab 0, Lab 1, ..." clearly at the top of your PDF.

2.1.3 Screenshot Annotations

On every screenshot containing SQL script, add a comment at the top with the following information:

1. Student Name
2. Department
3. Student Index Number

```
-- Name: student name  
-- Department: Computer Science, ...  
-- Index: 01x-xxx
```

2.1.4 Deadline

1. Submissions must be uploaded before the deadline.
2. Late submissions may be penalised according to the situations.

2.1.5 Plagiarism

- You may discuss the assignment with other students, but all work must be completed individually.
- Plagiarism will be identified and penalized appropriately based on the circumstances.
- The use of AI tools to complete your tasks is not allowed. If detected, it will result in penalties.

3. Setup environment

Objective:

- To set up **MySQL Database Management System**
- To connect to **MySQL** server
- To set up **Python Programming Language**


Software:

- MySQL installer (offline version)
- Visual Studio code (VS code)


3.1 Prerequisites: Install Microsoft Visual C++ 2019 Redistributable

If you already have Microsoft Visual C++ 2019 Redistributable installed, you can skip this step and continue with the lab.

1. Go to the official Microsoft page. [Download VC++ Redistributables](#)
2. Download the file based on your system:
 - for 64-bit [VC_redist.x64.exe](#)
 - for 32-bit [VC_redist.x86.exe](#)
3. Open the downloaded file and check on **I agree to the license terms and conditions**.



4. After installation, **restart your PC**.



3.2 MySQL

MySQL is an open-source relational database management system (RDBMS) that uses SQL to create, manage, and query data.

3.2.1 Why we will use MySQL?

1. Support **ACID transaction properites**, essential for understanding transaction management
2. Uses the **InnoDB** transactional storage engine, which supports **concurrency control** and **recovery mechanisms**
3. Flexible and easy-to-use

3.2.2 Installing MySQL (Offline)

1. Go to the official [MySQL Downloads Page](#)
2. Download the **Windows (x86, 32-bit), MSI Installer** (Offline version ~353.7M)

Note: MySQL 8.0 is the final series with MySQL Installer. As of MySQL 8.1, use a MySQL product's MSI or Zip archive for installation. MySQL Server 8.1 and higher also bundle MySQL Configurator, a tool that helps configure MySQL Server.


Select Version:
8.0.42

Select Operating System:
Microsoft Windows


Windows (x86, 32-bit), MSI Installer (mysql-installer-web-community-8.0.42.0.msi)	8.0.42	2.1M	Download
Windows (x86, 32-bit), MSI Installer (mysql-installer-community-8.0.42.0.msi)	8.0.42	353.7M	Download

We suggest that you use the [MD5 checksums](#) and [GnuPG signatures](#) to verify the integrity of the packages you download.

3. After open the Installer Choose **Custom**




4. Select **MySQL Servers** and **MySQL Workbench**




5. Proceed through the installation wizard

6. Use the default port (3306)




7. Choose the Authentication Method as below



8. Create a root password


9. Choose Windows Service as below




10. Proceed through the installation wizard until finish

3.2.3 What after install MySQL?


1. After installing MySQL, connect to the **MySQL Server**. by clicking the **+** symbol next to MySQL Connections



2. In the dialog box, enter a **Connection name**, e.g., `Local`




3. Click **Test Connection** and enter your MySQL root password to verify the connection



4. After successfully connecting to the MySQL server, create a new database called `StudentDB` :

- Open `local` connection
- Click **Create new SQL** from the top-left corner
- Write SQL commands to create the `StudentDB` database, a `Students` table, and insert sample data.
- Execute the Script using Execute button 



```

/*
Name: Amira Naser Aldeein
Index: 01x-xxx
Department: e.g., CS
*/

CREATE DATABASE IF NOT EXISTS StudentDB;
USE StudentDB;


CREATE TABLE Students (
  StudentID INT PRIMARY KEY AUTO_INCREMENT,
  FullName VARCHAR(100) NOT NULL,
  Gender ENUM('M', 'F') NOT NULL,
  BirthDate DATE,
  Department VARCHAR(100),
  GPA DECIMAL(3,2),
  CreatedAt TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Sample data
INSERT INTO Students (FullName, Gender, BirthDate, Department, GPA) VALUES
('Ali Ahmed', 'M', '2000-05-12', 'Computer Science', 3.45),
('Sara Hassan', 'F', '2001-07-18', 'Information Systems', 3.78),
('Omar Musa', 'M', '2002-03-25', 'Mathematics', 2.90);

```

 Once you create the database, you will see it listed under the Schemas section on the left panel

5. Now you can run SQL queries on your new database



```
USE StudentDB;
SELECT * FROM Students;
```

What do we mean by Connection and Port?

A database connection is a communication link between a client application (e.g., MySQL Workbench) and the MySQL server. When an application (the client) needs to perform operations such as creating databases, running queries, or retrieving data, it establishes a connection to communicate with the server. The port is a communication endpoint on your machine. By default, MySQL uses port 3306 to listen for incoming client requests. In short, a connection is essential to allow your client tools or applications to interact with the MySQL server, run queries, manage data, and exchange information.

3.3 Python programming language


We will use the **Python programming language** to simulate the following: * Concurrency control algorithms * Database recovery techniques

3.3.1 Download and Install Python


- 1. Go to the **Python** official website
- 2. Download the latest stable version
- 3. Open the installer and check
 - Add Python to PATH

3.3.2 Download and Install VS Code


- 1. Visit the website **VS code** Click on **Windows** to download VS code for Windows, as shown below



- 2. After the download finished open the `VSCoDeUserSetup` executable file, when it open select `I accept the agreement` then click on `Next`



3. Select all the option as seen below



4. Finally we are ready to install the VS code, click on `install` and wait until the setup finish

5. On the Vs code, on the left-hand side, click on `extension` then in the search bar, write `Python` select `Python microsoft` and click on `Install` as shown below



3.4 Alternative way to practice:

- Online Postgres server [supabase](#)
- Online Python Editor [programiz](#)

3.5 Assignment: Install MySQL, Create a Database, and Query It

 Due Date on 24/5/2025

See the requirement about the structures of the lab [here](#)

1. Download and Install MySQL
2. Connect to the server
3. Create a New Database
4. Create a Table
5. Insert Data
6. Query the Data

```
USE DatabaseName;
SELECT * FROM TableName;
```

3.5.1 What to assign:

Take Screenshots:

1. SQL code (on MySQL)
2. The output of your query
3. Put it all on one document

4. Transaction Management and SQL Implementation

Objective:

- Understand the ACID properties.
- Implement SQL transactions.
- Understand transaction control.

4.1 Understanding SQL Transactions

A **Transaction** is a sequence of read and write operations performed as a single logical unit of work.

SQL Server operates in the following transaction modes:

1. Autocommit transactions: Each individual statement is a transaction.
2. Explicit transactions: Each transaction is explicitly started with the `BEGIN TRANSACTION` statement and explicitly ended with a `COMMIT` or `ROLLBACK` statement (**which is concerned in this course**).
3. Implicit transactions: A new transaction is implicitly started when the prior transaction completes, but each transaction is explicitly completed with a `COMMIT` or `ROLLBACK` statement.

4.1.1 Transaction Control Commands – `SAVEPOINT`, `ROLLBACK`, and `COMMIT`

Objective: Understand how to manage multi-step database transactions using SQL transaction control commands.

BEGIN

Start a new transaction:

```
BEGIN;  
-- or  
START TRANSACTION;
```

SAVEPOINT

Set a point you can roll back to later:

```
BEGIN;  
  
INSERT INTO employees (id, name, department) VALUES (1, 'Alice', 'HR');  
SAVEPOINT sp1;  
  
INSERT INTO employees (id, name, department) VALUES (2, 'Bob', 'Finance');  
SAVEPOINT sp2;  
  
INSERT INTO employees (id, name, department) VALUES (3, 'Charlie', 'IT');
```

At this point, we have added 3 employees and created two savepoints: `sp1` and `sp2`.

ROLLBACK

Undo part or all of a transaction.

ROLLBACK TO A SAVEPOINT

```
ROLLBACK TO sp2;
```

This undoes only the insertion of **Charlie**, keeping **Alice** and **Bob**.

ROLLBACK ENTIRE TRANSACTION

```
ROLLBACK;
```

This undoes **all** changes made in the transaction.

COMMIT

Make all changes permanent:

```
COMMIT;
```

After committing, the changes cannot be undone.

Summary

- Use `SAVEPOINT` to mark logical checkpoints.
- Use `ROLLBACK` to undo mistakes.
- Use `COMMIT` to make sure the database saves your work.

This gives you control and safety when working with important or complex data operations.

4.1.2 Transaction properties and how SQL supports them

Transaction properties (ACID):

1. **A - Atomicity:** All actions in transaction happen, or none happen. “*All or nothing...*”
2. **C - Consistency:** If each transaction is consistent and the database starts in a consistent state, it ends in a consistent state. “*It looks correct to me...*”
3. **I - Isolation:** Execution of one transaction is isolated from other transactions. “*All by myself...*”
4. **D - Durability:** Once a transaction commits, its effects are permanent. “*I will survive...*”

How SQL Supports Each Property?

properties	SQL Feature
Atomicity	Use <code>START TRANSACTION</code> , followed by <code>COMMIT</code> or <code>ROLLBACK</code>
Consistency	Enforced through constraints (e.g., <code>PRIMARY KEY</code> , <code>FOREIGN KEY</code> , <code>CHECK</code>)
Isolation	Set using <code>SET TRANSACTION ISOLATION LEVEL</code>
Durability	Ensured by <code>COMMIT</code> and Logging mechanisms (e.g., write-ahead logs, redo logs)

4.1.3 Why Concurrency Control is Needed

Several problems can occur when concurrent transactions execute without proper control, especially due to conflicting operations such as **Read-Write**, **Write-Read**, and **Write-Write** conflicts.

Here is an illustrated problem with SQL transactions to simulate the problem:

1. Lost update problem

This problem occurs when two or more transactions read and update the same data item concurrently without being aware of each other's changes. As a result, the last update overwrites the previous ones, leading to the loss of some updates and causing data inconsistency.

Real-World Scenario

Two users (sessions) try to withdraw money from the same account at the same time:

- **User A** reads the balance = 100, decides to withdraw 30 → wants to update it to 70.
- **User B** also reads the balance = 100, decides to withdraw 50 → wants to update it to 50.
- If both updates happen without proper control, one will overwrite the other, leads to lost updates and incorrect final balance.


Example of the Scenario in SQL:

To simulate a Lost Update in MySQL using two concurrent sessions:

1. Create a Table: Create any Database and create a table named `accounts` and insert sample data into it:


```
CREATE TABLE accounts (
  id SERIAL PRIMARY KEY,
  name TEXT,
  balance NUMERIC(10,2)
);

INSERT INTO accounts (name, balance) VALUES ('Ali', 100.00);
```



2. Session A (user A):


```
START TRANSACTION;
SELECT balance FROM accounts WHERE id = 1;
UPDATE accounts SET balance = 70 WHERE id = 1;
```



Here, user A reads the balance (100) and calculates the new balance: $100 - 30 = 70$, then updates the record without commit.

3. Session B (user B, in another window):

```
START TRANSACTION;
SELECT balance FROM accounts WHERE id = 1;
UPDATE accounts SET balance = 50 WHERE id = 1;
COMMIT;
```




Here, user B reads the balance (100) and calculates the new balance:


$100 - 50 = 50$, then updates the record and commit.

to open two sessions in MySQL Workbench; Go to File → New Query Tab to open Session A, Repeat to open another tab (Session B). Each tab is a separate session — you can run SQL commands independently in each

4. Back to session A: Now return back to session A again and commit the session:



5. Check the final balance:



Here we will see the balance becomes 70, and user B's update (which should make the balance 50) is lost.

6. Conclusion:

- Both transactions read the same initial value (100).
 - Each transaction updates the data without being aware of the other's changes.
 - The last transaction to commit determines the final value, possibly erasing the other's update.
 - Even though no errors appear, data integrity is lost. This demonstrates the importance of proper concurrency control in database systems.
-

2. The Temporary Update (or Dirty Read) Problem

This occurs when one transaction reads data that has been written by another transaction that hasn't committed yet. If the first transaction is later **rolled back**, the second transaction ends up working with invalid or temporary data — a dirty read.

Example Scenario (Simulated):

In an online electronics store:

1. Ahmed initiates a purchase of 200 laptops. A transaction begins and temporarily reduces the stock quantity to 0, but the transaction is not yet committed.
2. Meanwhile, Mohamed checks the available stock and sees 0 laptops — based on User A's uncommitted change.
3. Then, Ahmed cancels the purchase, and the transaction is rolled back. The stock returns to its original quantity (e.g., 200).
4. However, Mohamed made a decision based on incorrect (temporary) data — the system gave them a dirty read.

SQL Transaction for the scienario:

1. Create the `Products` table to simulate the scenario

```
-- Create ProductDB Database
CREATE DATABASE IF NOT EXISTS ProductDB;
USE ProductDB;

-- Create Products table
CREATE TABLE Products (
  Id INT PRIMARY KEY,
  ProductName VARCHAR(100),
  Quantity INT
);

-- Insert Sample data
INSERT INTO Products (Id, ProductName, Quantity) VALUES
(1, 'Mobile', 100),
(2, 'Tablet', 50),
(3, 'Laptop', 200);
```

2. In Session 1: Connect using the previous connection from **Lab0** (named local), and open a new SQL query tab

```
USE ProductDB;

-- Transaction 1: Simulate purchase
START TRANSACTION;
-- Pay 200 units from product with ID = 3 (assumed to be a laptop)
UPDATE Products
SET Quantity = Quantity - 200
WHERE Id = 3;
-- Simulate waiting for payment confirmation
DO SLEEP(15);
-- Payment failed – rollback the transaction
ROLLBACK;
```


We use `SLEEP` to simulate concurrent transactions. Note that `DO SLEEP(15)` works only in MySQL 8.0.29 or later

3. Go to the Home tab and open the `local` connection again. (This step simulates a second session connected to the database.) Then, open a new SQL query tab to continue.

```
USE ProductDB;

-- Transaction 2: Retrieve Laptop record
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
START TRANSACTION;
SELECT *
FROM Products
WHERE Id=3;
COMMIT;
```


4. In Session 1: Run the transaction and wait




The screenshot shows the MySQL Workbench interface. The SQL editor contains the transaction script for Session 1. The Output tab at the bottom shows the execution progress of the transaction. The steps executed are:

#	Time	Action	Message	Duration / Fetch
1	16:04:30	USE ProductDB	0 row(s) affected	0.000 sec
2	16:04:30	START TRANSACTION	0 row(s) affected	0.000 sec
3	16:04:30	UPDATE Products SET Quantity = Quantity - 200 WHERE Id = 3	1 row(s) affected Rows matched: 1 Changed: 1 Warnings: 0	0.016 sec
4	16:04:30	DO SLEEP(15)	Running...	?

5. While Session 1 is still sleeping: Run Session 2



6. After Session 1 completes and rolls back



7. Conclusion: Session 2 was able to see the updated quantity (0 laptops) before Session 1 rolled back. This leads to inconsistent or misleading data, which is why isolation levels like `READ COMMITTED` or `REPEATABLE READ` are important in real-world applications.

3. The Incorrect Summary Problem

Problem Statement

When you use aggregate functions like `SUM`, `AVG`, and `COUNT` on a table while other transactions are modifying the data, the result may be **inconsistent**. This is because some rows might be updated **before** the function reads them, and others **after**.

1. Create the Table

```
CREATE TABLE accounts (
  id INT PRIMARY KEY,
  balance DECIMAL(10,2)
);
```

2. Insert Initial Data

```
INSERT INTO accounts (id, balance) VALUES
(1, 500.00),
(2, 750.00),
(3, 1200.00);
```

3. Simulate Concurrent Transactions

Transaction T1 (Session 1)

```
BEGIN;
UPDATE accounts SET balance = balance + 100 WHERE id = 1;
-- not committed
```

Transaction T2 (Session 2)

```
SELECT SUM(balance) FROM accounts;
```

At this point, `SUM()` still uses the old value for `id = 1` because the update in Session 1 (T1) hasn't been committed.

SUM(balance) = 2450.00

4. Add COMMIT to T1 and Re-run

Transaction T1 (Session 1)

```
BEGIN;
UPDATE accounts SET balance = balance + 100 WHERE id = 1;
COMMIT;
```

Transaction T2 (Session 2)

```
SELECT SUM(balance) FROM accounts;
```

Now, `SUM()` uses the new value for `id = 1` because the update in T1 has been committed.

SUM(balance) = 2550.00

5. Conclusion:

- If T1 is not committed, T2 sees old values.
- If T1 is committed, T2 sees updated values.
- Proper transaction management ensures consistent query results.

4.2 Assignment: Simulate The Unrepeatable Read Problem

 Due Date on 14/6/2025

- Read page 752 on the primary book (FUNDAMENTALS OF Database Systems)
- See the requirement about the structures of the lab [here](#)
- You can use any SQL editor (Online or Local), but be cautious — **transaction statements may vary slightly between different DBMSs**, and the use of `DO SLEEP` depends on the specific database engine. For example, `DO SLEEP()` is supported in MySQL 8.0.29+, while in other systems or older versions, you may need to use `SELECT SLEEP()` or an alternative approach.

4.2.1 What to assign:

1. A paragraph describing the scenario you will use to simulate the Unrepeatable Read problem.
2. A screenshot of the table you created for the simulation.
3. A screenshot showing the result of the first transaction after execution.
4. A screenshot showing the result of the second transaction after execution.
5. A paragraph discussing what happened, and how it demonstrates the Unrepeatable Read issue.

5. Concurrency Control Techniques

Objective:

- To understand concurrency control techniques based on locking mechanisms.
- To explore concurrency control techniques based on the multiversion (MVCC) concept.
- To understand how MySQL (InnoDB) handles concurrent transactions effectively.

5.1 InnoDB Storage Engine

A storage engine is a software that is used by a database management system to create, read, and update data from a database. **InnoDB** is a general purpose **transactional storage engine** for MySQL(since version 5.5.5), designed for high reliability and performance.

5.1.1 Why Do We Use InnoDB in This Course?

- Its DML operations follow the **ACID** model, supporting transactions with **commit**, **rollback**, and **crash recovery**, which help protect user data.
- It allows us to observe how **concurrency control algorithms** work within a DBMS, including **strict Two-Phase Locking (2PL)** and **Multi-Version Concurrency Control (MVCC)**.
- It enables us to examine **recovery mechanisms**, such as **write-ahead logging (WAL)**, in action.

5.1.2 How to check if InnoDB enabled on MySQL?

To verify whether InnoDB is enabled and set as the default storage engine, execute the following query:

```
SHOW ENGINES;
```

Look for the InnoDB row in the result. It should show **DEFAULT** in the Support column, indicating that InnoDB is the default storage engine.

InnoDB	DEFAULT	*Supports transactions, row-level locking, and foreign k...	YES	YES	YES
--------	---------	---	-----	-----	-----

5.1.3 InnoDB Locking

- **Shared and Exclusive Locks:** InnoDB implements standard row-level locking, which includes two types of locks: shared (S) locks and exclusive (X) locks.
- **Intention Locks** InnoDB supports multiple granularity locking, allowing row-level locks and table-level locks to coexist. It introduces two types of intention locks: intention shared (IS) and intention exclusive (IX). For details, refer to the lock compatibility matrix on page 803 of the book.

5.2 Isolation Level on MySQL

Isolation level ^{instruct} the database engine on how to manage multiple transactions being performed concurrently, and what violations are possible.

1. **READ UNCOMMITTED** is the lowest isolation level. It allows transactions to read the most recent version of a row, even if the change has not been committed by other transactions. This leads to the dirty read anomaly, as explained in [Lab 1](#).
2. **READ COMMITTED** is one step above **READ UNCOMMITTED** and prevents dirty reads. In this mode, each **SELECT** operation retrieves the latest committed version of the row at the time of the query. However, this can result in non-repeatable reads.
3. **REPEATABLE READ** is the default isolation level in MySQL. It is one step above **READ COMMITTED** and prevents non-repeatable reads. In this mode, shared locks are placed on all rows read during the transaction and are held until the transaction ends. This prevents other transactions from modifying any of the data that was read.
4. **SERIALIZABLE** is the highest isolation level. It prevents all types of concurrency anomalies. However, it is often impractical due to the extensive locking it requires, which increases the risk of deadlocks and can significantly reduce performance.

Syntax for Setting the Isolation Level in MySQL

```
SET TRANSACTION ISOLATION LEVEL [ISOLATION LEVEL];
```

 InnoDB use REPEATABLE READ as the default isolation level.

5.3 InnoDB-Supported Concurrency Control Algorithms

The InnoDB transaction model combines the strengths of Multi-Version Concurrency Control (MVCC) and strict Two-Phase Locking (2PL) to balance high concurrency with strong consistency guarantees. InnoDB uses MVCC to enable non-blocking reads, allowing multiple transactions to read data simultaneously without locking. At the same time, it applies strict 2PL for write operations to ensure serializability.

 When we refer to write operations, we mean SQL statements such as INSERT , UPDATE and DELETE

InnoDB uses two primary concurrency control mechanisms:

5.3.1 1. Two-Phase Locking (2PL)

InnoDB uses **strict two-phase locking** to ensure serializability. This means a transaction acquires locks during its execution phase but does not release any exclusive (write) locks until it either commits or rolls back

Example: Observing How InnoDB Uses Strict 2PL

Imagine Ahmed's family and Mohamed's family both trying to book seats for a summer vacation trip to Turkey. They access the reservation system at the same time. Ahmed wants to reserve five seats, while Mohamed wants to reserve four seats. Behind the scenes, InnoDB ensures data consistency by using strict two-phase locking (2PL) — meaning when Ahmed's transaction starts updating the available seat count, it locks the row exclusively. Mohamed's transaction, attempting to update the same row simultaneously, must wait until Ahmed's transaction commits or rolls back. This prevents overlapping writes and guarantees serializability.

Now we will simulate the example using MySQL and observe the InnoDB storage engine:

1. Create the Database and Table

```
CREATE DATABASE IF NOT EXISTS ReservationDB;
USE ReservationDB;

CREATE TABLE seats (
  Id INT PRIMARY KEY AUTO_INCREMENT,
  FlyName VARCHAR(100) NOT NULL,      -- Flight Name
  Available_seats INT                 -- Available seats on the flight
) ENGINE=InnoDB;

-- Insert some sample data
INSERT INTO seats (FlyName, Available_seats) VALUES
  ("Turkey", 10),
  ("Emaraty", 20);
```

2. In Session 1 (Ahmed's Transaction): Connect using the previous connection from **Lab0** (named local), and open a new SQL query tab

```
USE ReservationDB;

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
START TRANSACTION;

-- Ahmed tries to reserve 5 seats
UPDATE seats SET available_seats = available_seats - 5 WHERE id = 1;

-- Optional: Use SLEEP() to pause Ahmed's transaction.
-- This allows you to observe the lock behavior by giving Mohamed's session time to run concurrently.
-- If you prefer not to use SLEEP(), you can simply delay committing Ahmed's transaction
-- and manually run Mohamed's session before issuing the COMMIT.
DO SLEEP(20);

COMMIT;
```

3. In Session 2 (Mohamed's Transaction): Go to the Home tab and open the **local** connection again. (This step simulates a second session connected to the database.) Then, open a new SQL query tab to continue.

```
-- Mohamed's Transaction

USE ReservationDB;

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
START TRANSACTION;

-- Mohamed tries to reserve 4 seats
UPDATE seats SET available_seats = available_seats - 4 WHERE id = 1;


-- If Ahmed's transaction hasn't committed yet, this query will be BLOCKED.
-- It waits for the exclusive lock to be released.
COMMIT;
```

4. In Session 3

```
SHOW ENGINE INNODB STATUS;
```

`SHOW ENGINE INNODB STATUS` provides a detailed snapshot of the InnoDB storage engine's internal state. This output—often referred to as the InnoDB Monitor—includes several sections such as transaction and lock information, I/O activity, buffer pool usage, deadlock detection, and more.

- Run Session 1, Session 2, and Session 3 simultaneously. You will observe that the `UPDATE` operation in Session 2 keeps running (waiting) until Session 1 commits. Meanwhile, execute Session 3 during this wait period to view the InnoDB status and observe how Session 2 is blocked, waiting for Session 1 to release the lock.
- Right-click on the Status column in the InnoDB status output and select **Copy Row (with names, tab separated)**, then paste it into any text editor (e.g., Notepad) to examine the output in detail. In the snapshot below, you can see evidence that Transaction 2 is waiting for Transaction 1 to commit—this clearly demonstrates how InnoDB enforces strict two-phase locking (2PL).



Transaction 2

```
---TRANSACTION 4939, ACTIVE 3 sec starting index read
mysql tables in use 1, locked 1
LOCK WAIT 2 lock struct(s), heap size 1128, 1 row lock(s)
MySQL thread id 81, OS thread handle 14892, query id 3521 localhost 127.0.0.1 root updating
UPDATE seats SET available_seats = available_seats - 4 WHERE id = 1
----- TRX HAS BEEN WAITING 3 SEC FOR THIS LOCK TO BE GRANTED:
```

Transaction 1

```
---TRANSACTION 4938, ACTIVE 5 sec
2 lock struct(s), heap size 1128, 1 row lock(s), undo log entries 1
MySQL thread id 77, OS thread handle 14828, query id 3515 localhost 127.0.0.1 root User sleep
DO SLEEP(20)
```

Once the sleep duration completes, Transaction 1 commits and releases its exclusive lock, allowing Transaction 2 to acquire the lock and proceed. If you query the `available_seats`, you will find it equals 1, indicating that both transactions executed without any anomalies, as demonstrated in [Lab 1](#).

- Conculation on this example we observe that how InnoDB follow the strict 2PL on updating operation. you can repeat this example by using `READ COMMITTED`, `SERIALIZABLE` Isolation level and observe what happen.

5.3.2 2. Multiversion Concurrency Control (MVCC)

Definition

MVCC is a concurrency control technique used by many modern databases (like PostgreSQL, MySQL InnoDB, and Oracle) to allow multiple transactions to access the same data simultaneously without locking.

How MVCC Works

- Every time a transaction reads data, it sees a snapshot of the database at the time the transaction started.
- When data is updated, the database creates a new version of the row instead of overwriting the old one.
- Other transactions can continue reading the old version until they commit or refresh.

Objective

Understand how MVCC in InnoDB allows multiple read operations to happen at the same time, without locking, and while updates may be happening in the background.

Step 1: Create the Table

```
CREATE TABLE employees (  
  id INT PRIMARY KEY,  
  name VARCHAR(100),  
  salary INT  
) ENGINE=InnoDB;
```

Step 2: Insert Initial Data

```
INSERT INTO employees (id, name, salary) VALUES  
(1, 'Alice', 5000),  
(2, 'Bob', 6000);
```

Step 3: Start Two Sessions

We will simulate **Session A (Reader)** and **Session B (Updater)**. You can do this using two different database connections (e.g., two tabs in MySQL Workbench).

Session A (Reader)

```
START TRANSACTION;  
SELECT * FROM employees WHERE id = 2;
```

Observation: All rows are returning back, WHERE id = 2 returns 6000.

Session B (Updater) (In another tab)

```
START TRANSACTION;  
UPDATE employees SET salary = 9000 WHERE id = 2;  
COMMIT;
```

Back to Session A (Reader Again)

```
SELECT * FROM employees WHERE id = 2;
```

Observation: Now it shows 6000 for id=2.

Transaction A does not see the update from Transaction B because MVCC provides a consistent snapshot of the data at the moment A starts.

MVCC keeps multiple versions of rows, so Transaction A reads the version that existed before B's update. Even if B commits, A continues to see the old data until it finishes, ensuring repeatable reads without blocking.

5.3.3 Key Discussion Points

Feature	Observation
MVCC	Keeps a snapshot per transaction (using undo logs).
Multiple Reads	Readers don't block writers, and writers don't block readers.
Isolation	REPEATABLE READ is default in InnoDB – keeps results stable within a transaction.
Consistency	Session A always sees the same data until it commits.
Performance	Improves concurrency: multiple users can safely read and write.

5.3.4 Summary

This simple example shows how MVCC enables multiple users to read data at the same time, even if another transaction is updating the same rows. Each reader gets a consistent snapshot — a powerful feature of InnoDB.

5.4 Advanced work

You can clone this [repostory](#) to explore the internal details of concurrency control simulations.

5.5 Assignment (Simulate Deadlock Scenario)

 **Date on 28/6/2025**

- Read pages 789–792 from the primary textbook: Fundamentals of Database Systems.
- Please review the general lab structure requirements [here](#)
- You must use **MySQL**, as this lab focuses specifically on the InnoDB storage engine.

5.5.1 What to assign:

1. A brief paragraph describing the real-world situation you are simulating to create a deadlock.
2. A screenshot of the table structure (after creation) that you will use in the simulation.
3. A screenshot showing the result of the first session after execution.
4. A screenshot showing the result of the second session after execution.
5. While both Transaction 1 and Transaction 2 are still running and waiting on each other, open a third session and execute the following command: `SHOW ENGINE INNODB STATUS;`.
6. A snippet of the `SHOW ENGINE INNODB STATUS` output showing that Transaction 1 is waiting for a lock held by Transaction 2.
7. A snippet showing that Transaction 2 is waiting for a lock held by Transaction 1.
8. After the deadlock occurs and InnoDB resolves it automatically, re-run third session.
9. A snippet from the InnoDB Monitor output showing that a deadlock has been detected and resolved. to do this you should re-run step 5 again and observe the `LATEST DETECTED DEADLOCK`.
10. A brief paragraph discussing what happened and how InnoDB detects deadlock [Read this](#).